EunJeong Hwang

6724726353 hej78520@gmail.com
Website Google Scholar LinkedIn Github

Education

University of British Columbia **Sep. 2022 - Present**
PhD, Computer Science *4.0/4.0*

University of Massachusetts-Amherst **Sep. 2019 - May. 2021**
MS, Computer Science *3.67/4.0*

Kwangwoon University **Mar. 2014 - Feb. 2018**
BS, Computer Science *3.82/4.5*

Research Experience

University of British Columbia **Sep. 2022 - Present**
Research Assistant, Advised by Vered Shwartz *Vancouver, BC*

- Research about understanding implicit cues (such as user opinions, facial expressions, gestures, and beliefs) within explicit contexts, and modeling how to control the knowledge that AI systems rely on.
- Current project investigates **how to make mental simulation of LLMs more efficient during multi-turn conversation**.

Google DeepMind **Jan. 2024 - Jun. 2024**
Research Intern (Host: Yichao Zhou, Beliz Gunel) *Remote*

- Research on incremental structured summaries and resolving knowledge conflicts.
- Two papers (SUMIE: A Synthetic Benchmark for Incremental Entity Summarization & Enhancing Incremental Summarization with Structured Representations) accepted at EMNLP 2024 and COLING 2025.

MIT-IBM Research **Jul. 2023 - Nov.2023**
Research Intern (Host: Veronika Thost, Dan Gutfreund) *Boston, MA*

- Research on reasoning subjective opinions using graph structures.
- 'A Graph per Persona: Reasoning about Subjective Natural Language Descriptions' accepted at ACL-Findings 2024.

Allen Institute for AI (AI2) **Apr. 2023 - Dec.2023**
Research Collaborator (Mentor: Bodhisattwa Prasad Majumder, Niket Tandon) *Remote*

- Research on aligning language models to user opinions.
- 'Aligning Language Models to User Opinions' accepted at EMNLP-Findings 2023.

UMass Amherst-IBM **Feb. 2021 - Oct.2021**
Graduate Student Researcher (Mentor: Veronika Thost, Tengfei Ma) *Remote*

- Research on virtual node augmented graph neural networks (GNN) for link prediction tasks.
- 'An Analysis of Virtual Nodes in Graph Neural Networks for Link Prediction (Extended Abstract)' accepted as a spotlight at LoG 2022.

Information Extraction & Synthesis Lab (IESL), UMass Amherst **Jan. 2020 - Aug.2021**
Graduate Student Researcher (Mentor: Jay-Yoon Lee, Michael Boratko) *Amherst, MA*

- Research on box embeddings for relation extraction and question-answering tasks.
- 'Event-Event Relation Extraction using Probabilistic Box Embedding' accepted at ACL 2022.

Publications **(* denotes equal contribution)**

Agents of Chaos **Arxiv Preprint**

Natalie Shapira, Chris Wendler, ..., **EunJeong Hwang**, ..., Tomer Ullman, David Bau

My contribution: First to identify and wrote **"Owner Identity Spoofing"** vulnerability (Chapter 8) and wrote the Related Work section on **Theory of Mind limitations in agentic systems**. See Interaction logs here.

Infusing Theory of Mind into Socially Intelligent LLM Agents **Submitted to ACL 2026**

EunJeong Hwang*, Yuwei Yin*, Giuseppe Carenini, Peter West, Vered Shwartz

Watch Before You Answer: Learning from Visually Grounded Post-Training **Submitted to ECCV 2026**
Yuxuan Zhang, **EunJeong Hwang**, Yiming Jia, Penghui Du, Peter West, Kelsey R Allen

SWI: Speaking with Intent in Large Language Models **INLG 2025**
Yuwei Yin, **EunJeong Hwang**, Giuseppe Carenini

BottleHumor: Self-Informed Humor Explanation using the Information Bottleneck Principle **ACL-Findings 2025**
EunJeong Hwang, Peter West, Vered Shwartz

SUMIE: A Synthetic Benchmark for Incremental Entity Summarization **COLING 2025**
EunJeong Hwang*, Yichao Zhou*, Beliz Gunel, James Bradley Wendt, Sandeep Tata

From Local Concepts to Universals: Evaluating the Multicultural Understanding of Vision-Language Models **EMNLP 2024**
Mehar Bhatia, Sahithya Ravi, Aditya Chinchure, **Eunjeong Hwang**, Vered Shwartz

Enhancing Incremental Summarization with Structured Representations **EMNLP-Findings 2024**
EunJeong Hwang, Yichao Zhou, James Bradley Wendt, Beliz Gunel, Nguyen Vo, Jing Xie, Sandeep Tata

A Graph per Persona: Reasoning about Subjective Natural Language Descriptions **ACL-Findings 2024**
EunJeong Hwang, Vered Shwartz, Dan Gutfreund, Veronika Thost

MemeCap: A Dataset for Captioning and Interpreting Memes **EMNLP 2023**
EunJeong Hwang, Vered Shwartz

Aligning Language Models to User Opinions **EMNLP-Findings 2023**
EunJeong Hwang, Bodhisattwa Prasad Majumder, Niket Tandon

Knowledge Graph Compression Enhances Diverse Commonsense Generation **EMNLP 2023**
EunJeong Hwang, Veronika Thost, Vered Shwartz, Tengfei Ma

An Analysis of Virtual Nodes in Graph Neural Networks for Link Prediction (Extended Abstract) **LoG 2022 (Oral)**
EunJeong Hwang, Veronika Thost, Shib Sankar Dasgupta, Tengfei Ma

Event-Event Relation Extraction using Probabilistic Box Embedding **ACL 2022**
EunJeong Hwang, Jay-yoon Lee, Tianyi Yang, Dhruvesh Patel, Dongxu Zhang, Andrew McCallum

Work Experience

Naver Corp. **Mar. 2021 - Jun. 2022**
Software Engineer in Knowledge Base team *Seongnam, Korea*

- Constructed a knowledge graph based on regional district information, such as restaurants, festivals, and exhibitions, and expanded search queries using neo4j along with developing a weather-related search system.

IBM **Jan. 2018 - Apr. 2019**
Application Developer *Seoul, Korea*

- Stock Market Predictions using AI, Shinhan Financial Group: Extracted financial text patterns using a rule-based Watson solution and constructed a financial keyword graph by extracting important phrases in retrieved news.
- Information System Audit, Standard Chartered Bank: Maintained and improved functionality of bank surveillance system and bank branch audit system.
- Learning Platform Development, Hanwha: Implemented basic functions on the learning platform, including search, paging, listing, and export, and participated in all phases of the project from requirement specifications to user acceptance testing with clients.

IBM **Sep. 2017 - Dec. 2017**
Application Developer Intern *Seoul, Korea*

- Designed database tables and developed the learning platform websites in accordance with user interface design.

Invited Talks

Invited to give a talk at Boson AI (Host: Alex Smola (CEO of BosonAI)) **Happening soon in the summer of 2026!**
- Talk about my research.

Invited talk at ResearchTrend.ai (Host: Prof. SeongJoon Oh) **Nov. 26, 2025**
- Talk about 'Infusing Theory of Mind into Socially Intelligent LLMs'. (youtube link)

Invited talk at Seoul National University (Host: Prof. Jay-yoon Lee) **Sep 9, 2024**
- Talk about 'Can MLLMs reason on extra-linguistic contexts?'.

Career Panel Discussion at UMass-Amherst (Host: Prof. Andrew McCallum at 696DS) **Apr. 2024, May 2025**
- Discussion about the career choice between Industry and PhD program for Master's students.

Invited talk at UT Austin (Host: Prof. Jessy Li at LIN 393) **Feb 8, 2024**
- Talk about Aligning language models to user opinions.

Teaching Experience

TA for CPSC 532V Commonsense Reasoning (Prof. Vered Shwartz) **Jan. 2026 - Apr. 2026**
- Design, help and grade students' assignments and projects.

TA for CPSC 532V Commonsense Reasoning (Prof. Vered Shwartz) **Jan. 2025 - Apr. 2025**
- Help and grade students' assignments and projects.

Guest lecture about Multimodal Commonsense in CPSC 532V Commonsense Reasoning (Prof. Vered Shwartz) **Feb. 2024**
- Lectured about tasks requiring commonsense beyond textual modalities, recent vision-language models, and their future directions.

Grader for CS685 Advanced Natural Language Processing (Prof. Brendan O'Connor) **Feb. 2021 - May. 2021**
- Grade students' assignments, provide feedback, and help them out to solve confusion on class material.

Services

Workshop Co-Organizer
- SocialLLM: Large Language Models for Social Reasoning and Simulation @ ICWSM 2026
- Future of NLP Workshop at UBC (pre-NeurIPS event)
- 1st Personalization of Generative AI workshop (PERSONALIZE @ EACL 2024)

Reviewer/Program Committee
- ECCV(2026), ARR (2024, 2025, 2026 -), COLM (2025, 2026), CVPR (2025), ICLR (2025), CIKM (2024, 2025), Actionable Interpretability Workshop (AIW) at ICML 2025, *SEM 2024, The First Workshop on Data Contamination (CONDA 2024), AAAI Workshop on Responsible Language Models (ReLM 2024)

Other Service
- Volunteer at The Women in Machine Learning Symposium (WiML @ ICML 2025)
- Vector Institute Scholarship Adjudication Committee (2025, 2026)
- UBC CS: Admission Committee Member for the CSE Ph.D. Program (2022)

Mentoring
- Mentor and judge @ cmd-f: Western Canada's largest hackathon for underrepresented genders in tech (2025)

Awards

Travel Grant — IVADO **Mar. 2025**
- Awarded $3000 CAD to attend the week-long seminar, Safety & Guaranteed LLMs, at the Simons Institute, UC Berkeley

Fully Funded PhD — University of British Columbia **Sep. 2022 – Present**
- President's Academic Excellence Initiative PhD Award
- International Tuition Award
- Faculty of Science PhD Tuition Award

Merit Scholarship — KwangWoon University **Mar. 2014 – Sep. 2015**
- Received 3 Academic Excellence Scholarships (each around $2000 CAD) for outstanding GPA

Technical Skills

Programming: Python, Java
Packages: PyTorch

Hobbies

Kick boxing, Scuba diving, Tennis, Hiking